

Yukai Zhou (周宇凯)

[Personal Website](#)
[Google Scholar](#)
[frank-thou](#)
[zhouyk12023@shanghaitech.edu.cn](#)

Education

PhD	ShanghaiTech University, Computer Science	Sep. 2025 - Present
	- Advised by Prof. Wenjie Wang at ASPIRE Lab. - M.S. phase of the integrated M.S.–Ph.D. program from 2023 to 2025.	
BS	ShanghaiTech University, Physics	Sep. 2019 - Jun. 2023

Publications & Preprints

How Jailbreak Attacks Inform Safety Alignment: A Defender-Centric, Shapley-Based Evaluation of Jailbreak Contributions	Jan. 2026 - Sep. 2026
Yukai Zhou , Feiyang Lu, Xiaokai Mao, Jinfei Liu, Wenjie Wang [PDF] [Code] [Website]	
- Reframe jailbreak attacks as red-teaming resources and evaluate them by the downstream safety gains they enable, rather than attack success rate alone. - Propose a Shapley value estimation method that achieves a better accuracy–efficiency trade-off than baseline methods in controlled experiments.	
Beyond Jailbreaks: Revealing Stealthier and Broader LLM Security Risks Stemming from Alignment Failures	Feb. 2025 - Sep. 2026
Yukai Zhou , Sibe Yang, Wenjie Wang [PDF] [Code] [Website] [Dataset]	
- Identify one unexplored realworld LLM safety issue, and construct the first dataset & attack methods within this domain. - Conduct large-scale human study ($n>600$) by the prolific in the future revision.	
Don't Say No: Jailbreaking LLM by Suppressing Refusal	Jan. 2024 - Feb. 2025
Yukai Zhou , Jian Lou, Zhijie Huang, Zhan Qin, Sibe Yang, Wenjie Wang [PDF] [Code] [Chinese Blog] [Leaderboard of JailbreakBench]	
- Propose one novel and powerful learning-based attack. Address several existing issues such as "Loss-ASR Mismatch" Problem and trustworthy evaluation metric. - Ranked the 1-st in the JailbreakBench white-box leaderboard, with 120+ google scholar citations.	

Invited Talks

Beyond Attack Success Rate: Characterizing LLM Jailbreak Risks (upcoming)	Oct. 2026
Invited oral presentation at the inaugural “Jiao-Chuang” Global Doctoral Student Academic Forum, Shanghai Jiao Tong University. Forum notice	

Patent Applications

One filed patent application , related to the AttackSHAP work.	Jun. 2026
---	-----------

Projects and Contests

JailbreakBench , secure the first place in the white-box attack leaderboard	Sep. 2024 - Oct. 2024
The Competition for LLM and Agent Safety 2024 NeurIPS 2024 Contest	Sep. 2024 - Oct. 2024
Secure the best white-box method, and overall the second best* (before re-evaluation).	

CCF LLM safety challenge contest track 1: General purpose LLM hijack Contest website , three-person team from ASPIRE Lab	Jun. 2024 - Aug. 2024
Design and implement the LLM hijack optimization-based workflow and dataset.	
KG4SLvis: A Visual Analytics Approach to Illuminate KG4SL Predictions CS286 course project within a four-person group	Nov. 2023 - Jan. 2024
- Design and implement the full-stack visual analytics system to enhance the interpretability and performance of Synthetic Lethality prediction model (KG4SL). - Refine the KG4SL model using insights got from the system, conduct experiments to validate its effectiveness and robustness.	
Visual Prompt Tuning Optimization Investigate the initialization methods regarding the visual prompts, which falls into the Parameter-Efficient Fine-Tuning (PEFT) methods category.	Sep. 2023 - Nov. 2023

Award, Service, and Additional Experience

Reviewer , ACL ARR 2025, ACL ARR 2026, ACM Computing Surveys, ICLR 2027, etc.	2025 - 2026
Guest Lecturer , to give a lecture upon jailbreak in course CS246 (Trustworthy ML)	2025.4.2, 2026.4.15
Outstanding Student , awarded to the top 10%	Sep. 2023 - Aug. 2024
Teaching Assistant , Introduction to Information Science and Technology (SI100b)	Mar. 2024 - Jun. 2024
Mentor of Dadao college , Mentoring the undergraduate students of Dadao college	Sep. 2023 - Jan. 2024